

Sample Doc 1: RAG Basics

Retrieval-Augmented Generation (RAG) combines a retriever and a generator. The retriever fetches evidence chunks from documents.

A common baseline is TF-IDF retrieval. Another baseline is BM25, which uses term frequency and inverse document frequency.

Good RAG answers should be grounded in the retrieved evidence and should not hallucinate facts that are not supported.

When evidence is missing, the system should say 'I don't know' or request more context.